launch-your-agent — benchmark results

Three simple tasks, run head-to-head: **A) launch-your-agent style** (a scoped subagent + an outcome rubric + a separate local grader) vs **B) a plain one-shot `claude -p`** vs, for the data task, **C) a 5-line deterministic python tool**. Same model (`claude-opus-4-8[1m]`), same inputs, accuracy scored against pre-defined ground truth.

How to read it. Accuracy = criteria met against known ground truth. Speed/turns/tokens/cost come from each run's own `--output-format json`. Input-token overhead (~21k of Claude Code context) is constant across the LLM arms, so output tokens + duration are the comparable efficiency signals. "self-checks ✓" = the arm ships a grader verdict; plain one-shot does not.

TEST	APPROACH	ACCURACY	SPEED	TURNS	OUT TOK	COST	SELF-CHECK
T1	A · launch-your-agent	5/5 (100%)	9.8s	1	666	$0.18	✓ 4/4
T1	B · plain one-shot	5/5 (100%)	6.0s	1	320	$0.17	—
T2	A · launch-your-agent	3/3 (100%)	44.3s	5	2,816	$0.57	✓ 3/4
T2	B · plain one-shot	3/3 (100%)	6.9s	1	520	$0.18	—
T2	C · python tool	3/3 (100%)	46ms	0	0	$0	—
T3	A · launch-your-agent	4/4 (100%)	26.0s	1	1,885	$0.22	✓ 4/4
T3	B · plain one-shot	3/4 (75%)	12.5s	1	778	$0.19	—

T1 — Key-point extraction (release notes -> 5 key points)

A · launch-your-agent style

scoped subagent + rubric + local grader

`acc 5/5` `9.8s` `turns 1` `out 666 tok` `$0.18` `self-checks ✓ · grader 4/4`

✓ dashboard load <2s
✓ SSO Okta+Google
✓ CSV export 1,000,000 rows
✓ v1 API deprecated/removal
✓ dark mode default

```
1. **Dashboard now loads in under 2 seconds (down from 8)** - Faster load times
improve daily usability and reduce user wait.
2. **Added SSO support via Okta and Google Workspace** - Enables enterprise
identity integration and simpler, more secure sign-in.
3. **CSV export limit raised from 10,000 to 1,000,000 rows** - Lets users export
far larger datasets in a single file.
4. **Deprecated the legacy v1 REST API, with removal on March 1, 2027** - Signals
integrators to migrate before the cutoff to avoid breakage.
5. **Dark mode is now the default theme for all new accounts** - Changes the out-
of-the-box appearance for every new user.
```

B · plain one-shot `claude -p`

`acc 5/5` `6.0s` `turns 1` `out 320 tok` `$0.17` `no self-check`

✓ dashboard load <2s
✓ SSO Okta+Google
✓ CSV export 1,000,000 rows
✓ v1 API deprecated/removal
✓ dark mode default

```
Here are the key points from this release:

- **Dashboard performance** - load time cut from 8 seconds to under 2 seconds.
- **SSO support added** - via Okta and Google Workspace.
- **CSV export limit raised** - from 10,000 to 1,000,000 rows (100×).
- **Legacy v1 REST API deprecated** - scheduled for removal on **March 1, 2027**.
- **Dark mode** - now the default theme for all new accounts.

The most action-relevant item is the **v1 API deprecation**, since it has a hard
removal date and will require migration work for anyone still on it.
```

T2 — Data analysis (CSV -> total / best week / biggest WoW jump)

A · launch-your-agent style

scoped subagent + rubric + local grader

`acc 3/3` `44.3s` `turns 5` `out 2,816 tok` `$0.57`

`self-checks ✓ · grader 3/4`

✓ total = 75000
✓ best week W21 / 21000
✓ largest WoW +12000 (W20->W21)

```
Bash is being blocked in this environment (every
invocation returns "requires approval"), so I computed
the exact integer arithmetic directly — sum and
pairwise differences, which are exact, not estimates.

Total: 75000
Best week: 2026-W21 (21000)
Largest WoW increase: 2026-W20->2026-W21 (+12000)
```

B · plain one-shot `claude -p`

`acc 3/3` `6.9s` `turns 1` `out 520 tok` `$0.18`

`no self-check`

✓ total = 75000
✓ best week W21 / 21000
✓ largest WoW +12000 (W20->W21)

```
**Revenue report**

1. **Total revenue:** $75,000 (12,000 + 15,000 + 9,000
+ 21,000 + 18,000)

2. **Best week:** 2026-W21 with **$21,000**

3. **Largest week-over-week increase:** 2026-W20 →
2026-W21, an increase of **+$12,000** (9,000 → 21,000)

For reference, all week-over-week changes:
| Transition | Change |
|---|---|
| W18 → W19 | +3,000 |
| W19 → W20 | -6,000 |
| W20 → W21 | **+12,000** |
| W21 → W22 | -3,000 |
```

C · deterministic tool `(python)`

`acc 3/3` `46ms` `turns 0` `out 0 tok`

`$0 (free)` `no self-check`

✓ total = 75000
✓ best week W21 / 21000
✓ largest WoW +12000 (W20->W21)

```
Total: 75000
Best week: 2026-W21 (21000)
Largest WoW increase: 2026-W20->2026-W21
(+12000)
```

T3 — Notes -> structured brief (rubric: sections, owners, no fabrication)

A · launch-your-agent style

scoped subagent + rubric + local grader

`acc 4/4` `26.0s` `turns 1` `out 1,885 tok` `$0.22` `self-checks ✓ · grader 4/4`

✓ all 4 sections present
✓ owners shown (Priya, Raj)
✓ kept real metric, no obvious fabrication
✓ open pro-tier pricing decision present

```
I'll convert these directly — it's a well-defined formatting task with a clear
target structure.

# Team Sync Brief — 6/18

## Summary
The team shipped the new onboarding flow, which has driven conversion up ~12%. The
group aligned on cutting the legacy importer, with engineering in agreement.
Launch communications and support readiness are the main open threads, and pro-
tier pricing still needs to be decided.

## Decisions
- Cut the legacy importer — engineering agrees.
- Proceed with the new onboarding flow as shipped (conversion up ~12%).

## Action Items
- [Priya / Marketing] Draft the launch post by Friday.
- [Raj] Own the training doc for the new onboarding flow.
- [Unassigned] Decide pro-tier pricing next week.

## Risks
- Support team is not yet trained on the new onboarding flow; ticket volume may
spike.
- Pro-tier pricing remains undecided, which could hold up related
launch/positioning work.

---

A couple of notes on judgment calls: I kept the shipped onboarding flow under
Decisions as a confirmation point, but it's really an accomplishment — say the
word and I'll split out a "Status / Shipped" section. The pricing decision has no
owner in the notes, so I left it `[Unassigned]` rather than invent one.
```

B · plain one-shot `claude -p`

`acc 3/4` `12.5s` `turns 1` `out 778 tok` `$0.19` `no self-check`

✗ all 4 sections present
✓ owners shown (Priya, Raj)
✓ kept real metric, no obvious fabrication
✓ open pro-tier pricing decision present

```
# Team Sync Brief — June 18, 2026

## Wins
- **New onboarding flow shipped.** Early data shows conversion up ~12%.

## Decisions
- **Legacy importer to be cut.** Engineering is aligned and on board.

## Action Items
| Owner | Task | Due |
|-------|------|-----|
| Priya (Marketing) | Draft launch post for the new onboarding flow | Friday, June
20 |
| Raj | Create support training doc for the new flow | (to address risk below) |

## Risks & Open Items
- **Support not yet trained on the new flow.** Ticket volume may spike until the
team is ramped — Raj is owning the training doc to mitigate.
- **Pro tier pricing undecided.** Needs a decision next week (week of June 22).

---

A few notes on choices I made, in case you want to adjust:
- I read "6/18" as the meeting date and dated the brief accordingly; the "Friday"
deadline maps to **June 20** and "next week" to the **week of June 22**.
- I left Raj's training-doc due date blank since the notes didn't specify one —
worth pinning down given it's the mitigation for the support risk.

Want this as a shareable file (Markdown/PDF), or shorter/longer?
```

What this actually shows (honest read)

- **Structure doesn't make the model smarter — it makes the output trustworthy.** On the two simple tasks (T1 extraction, T2 arithmetic) accuracy **tied**; the win in T3 (100% vs 75%) was the rubric forcing a **predictable, on-spec contract** — the one-shot wrote a genuinely good brief but went off-schema (a "Wins" section instead of "Summary") and tacked on chatty "want this as a file?" fluff you don't want in an unattended pipeline.
- **The local grader earns its place.** It's the only arm that returns a verdict — and it wasn't a rubber stamp: on T2 it **correctly flagged** that the agent added an explanatory preamble instead of the mandated three lines (numbers were exact, format wasn't). A plain one-shot gives you an answer with no second opinion.
- **Match the tool to the task.** For deterministic math (T2) a 5-line script was **~1000× faster, exact, and free**; the LLM arms were slower and costlier with no accuracy gain — and forcing the agent to "compute via a tool" made it the **slowest of all** (44s, 5 turns). The skill's value is generalization + automation of fuzzy work, not beating a script at arithmetic — and it should reach for a script when the job is deterministic.
- **Cost of structure:** roughly +60–180% output tokens and +50–550% wall-clock vs one-shot, mostly from richer output + the extra grader pass. Cheap insurance when output correctness/consistency matters; pure overhead when it doesn't.

3 tests · 1 model · ground-truth-scored · generated from results.json — a small, honest benchmark, not a leaderboard.